


# 1 Linear regression (for regression)

Model a scalar target with one or more quantitative features.

Although regression computes a linear combination, features can be transformed by nonlinear functions if relationships are known or can be guessed.



## 1.1 Univariate linear regression

- Description

Simple linear regression is a statistical method that studies the relationship between two variables:

- $x$ : the predictor, explanatory, independent variable,
- $y$ : the response, outcome, dependent variable.

- Model's hypothesis

$$h(x) = \theta_0 + \theta_1 * x = \hat{y}$$

- Model's parameters (# 2)

- $\theta_0, \theta_1$

- Cost function (in this case, the squared error function)

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (\hat{y}^i - y^{(i)})^2$$

- Goal

$$\text{minimize}_{\theta_0, \theta_1} J(\theta_0, \theta_1)$$

- Algorithm

- gradient descent

- start with some initial values for  $\theta_0$  and  $\theta_1$  (usually zero)

- keep changing  $\theta_0$  and  $\theta_1$  to reduce  $J(\theta_0, \theta_1)$

- $\theta_0 = \theta_0 - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1) = \theta_0 - \alpha \left( \frac{1}{m} \sum_{i=1}^m (\hat{y}^i - y^{(i)}) \right)$

- $\theta_1 = \theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1) = \theta_1 - \alpha \left( \frac{1}{m} \sum_{i=1}^m (\hat{y}^i - y^{(i)}) \cdot x^{(i)} \right)$

- directly use the following **formulas**

- $\theta_1 = \frac{\sum_{i=1}^m (x^{(i)} - \bar{x})(y^{(i)} - \bar{y})}{\sum_{i=1}^m (x^{(i)} - \bar{x})^2}$

- $\theta_0 = \bar{y} - \theta_1 \bar{x}$

- Notations

- $x$ : the independent variable
- $y$ : the dependent variable
- $m$ : the number of training examples
- $x^{(i)}$ : the input value of the  $i^{th}$  training example
- $y^{(i)}$ : the target value of the  $i^{th}$  training example
- $\hat{y}^{(i)}$ : the prediction made on the  $i^{th}$  training example by the current hypothesis function
- $\alpha$ : the learning rate; determines how big steps we take when updating the  $\theta$  parameters

- Hyperparameters:
  - $\alpha$ : if too small, slow gradient descent; if too large, gradient descent may fail to converge
- Problems:
  - gradient descent may converge to a local minimum

## 1.2 Multivariate linear regression

- Description

Multivariate linear regression is a statistical method that studies the relationship between multiple variables:

- $n$  variables  $x = \{x_1, x_2, \dots, x_n\}$ : the predictor, explanatory, independent variables,
- one  $y$  variable: the response, outcome, dependent variable.

- Model's hypothesis

$$h(x) = \theta_0 + \theta_1 * x_1 + \theta_2 * x_2 + \dots + \theta_n * x_n = \hat{y}$$

- Model's parameters (#  $n + 1$ )

$$\theta = \{\theta_0, \theta_1, \dots, \theta_n\}$$

- Cost function (in this case, the squared error function)

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)})^2$$

- Goal

$$\text{minimize}_{\theta} J(\theta)$$

- Algorithm

- gradient descent
  - start with some initial values for  $\theta_0, \theta_1, \dots, \theta_n$  (usually 0)
  - keep changing  $\theta$ s to reduce  $J(\theta)$ 
    - $\theta_0 = \theta_0 - \alpha \frac{\partial J}{\partial \theta_0} = \theta_0 - \alpha \left( \frac{1}{m} \sum_{i=1}^m (\hat{y}^i - y^{(i)}) \right)$
    - $\theta_1 = \theta_1 - \alpha \frac{\partial J}{\partial \theta_1} = \theta_1 - \alpha \left( \frac{1}{m} \sum_{i=1}^m (\hat{y}^i - y^{(i)}) \cdot x_1^{(i)} \right)$
    - ...
    - $\theta_n = \theta_n - \alpha \frac{\partial J}{\partial \theta_n} = \theta_n - \alpha \left( \frac{1}{m} \sum_{i=1}^m (\hat{y}^i - y^{(i)}) \cdot x_n^{(i)} \right)$

- Notations

- $x$ : the independent variables
- $y$ : the dependent variable
- $m$ : the number of training examples
- $n$ : the number of features representing each training example
- $x^{(i)}$ : the input values of the  $i^{th}$  training example
- $x_j^{(i)}$ : the value of the  $j^{th}$  feature of the  $i^{th}$  training example
- $y^{(i)}$ : the target value of the  $i^{th}$  training example
- $\hat{y}^{(i)}$ : the prediction made on the  $i^{th}$  training example by the current hypothesis function
- $\alpha$ : the learning rate; determines how big steps we take when updating the  $\theta$  parameters


- Hyperparameters:

- $\alpha$

- Problems:

- make sure features are on similar scales; gradient descent may be slow otherwise

- rescaling:  $x' = \frac{x - \min(x)}{\max(x) - \min(x)}$
- mean normalization:  $x' = \frac{x - \text{mean}(x)}{\max(x) - \min(x)}$
- standardization:  $x' = \frac{x - \text{mean}(x)}{\text{std}(x)}$



- make sure the gradient descent is working correctly

- plot the value of the cost function  $J$  over the number of iterations (# *epochs*)
- for a sufficiently small  $\alpha$ ,  $J(\theta)$  should decrease on every iteration
- if  $\alpha$  is too small, the gradient descent can be slow to converge
- if  $\alpha$  is too large,  $J(\theta)$  may not decrease on every iteration; may not converge
- try values  $\alpha \in \{\dots, 0.001, 0.01, 0.1, 1, \dots\}$

- in case of **underfitting**

- try adding new features
  - e.g. use a polynomial regression

$$\theta_0 + \theta_1 * x \Rightarrow \theta_0 + \theta_1 * x + \theta_2 * x^2 + \theta_3 * x^3$$

- in case of **overfitting**

- get more training examples
- try reducing the number of features
  - manually select which features to keep
  - use a model-selection algorithm
- try getting additional features
- try adding polynomial features
- use regularization
  - keep all the features, but reduce the magnitude of parameters  $\theta$ ; works well when working with a lot of features, each of which contributes a bit to predicting  $y$
  - intuition: *shrink* model parameters in order to *smooth out* the decision boundary (generate a *simpler* hypothesis)
  - cost function with the *regularization term*

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)})^2 + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

- parameter update in gradient descent (for  $j \in \{1, 2, \dots, n\}$ )

$$\theta_j = \theta_j - \alpha \frac{\partial J}{\partial \theta_j} = \theta_j - \alpha \left( \frac{1}{m} \sum_{i=1}^m (\hat{y}^i - y^{(i)}) \cdot x_j^{(i)} + \frac{\lambda}{m} \theta_j \right) = \theta_j \left( 1 - \alpha \frac{\lambda}{m} \right) - \alpha \frac{1}{m} \sum_{i=1}^m (\hat{y}^i - y^{(i)}) \cdot x_j^{(i)}$$

- $\lambda$  is the *regularization parameter* and needs to be tuned; it controls the trade-off between the goal of fitting the training data well and the goal of keeping the parameters small
- when using regularization, try increasing or decreasing  $\lambda$