


1 Linear regression (for regression)

Model a scalar target with one or more quantitative features.

Although regression computes a linear combination, features can be transformed by nonlinear functions if relationships are known or can be guessed.



Example of prediction using a linear regression model

1.1 Univariate linear regression

- Description

Simple linear regression is a statistical method that studies the relationship between two variables:

- x : the predictor, explanatory, independent variable,
- y : the response, outcome, dependent variable.

- Model's hypothesis

$$h(x) = \theta_0 + \theta_1 * x = \hat{y}$$

- Model's parameters (# 2)

- θ_0
- θ_1

- Cost function (in this case, the squared error function)

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (\hat{y}^i - y^{(i)})^2$$

- Goal

$$\text{minimize}_{\theta_0, \theta_1} J(\theta_0, \theta_1)$$

- Algorithm

- gradient descent

- start with some initial values for θ_0 and θ_1 (usually zero)

- keep changing θ_0 and θ_1 to reduce $J(\theta_0, \theta_1)$

- $\theta_0 = \theta_0 - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1) = \theta_0 - \alpha \left(\frac{1}{m} \sum_{i=1}^m (\hat{y}^i - y^{(i)}) \right)$

- $\theta_1 = \theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1) = \theta_1 - \alpha \left(\frac{1}{m} \sum_{i=1}^m (\hat{y}^i - y^{(i)}) \cdot x^{(i)} \right)$

- directly use the following **formulas**

- $\theta_1 = \frac{\sum_{i=1}^m (x^{(i)} - \bar{x})(y^{(i)} - \bar{y})}{\sum_{i=1}^m (x^{(i)} - \bar{x})^2}$

- $\theta_0 = \bar{y} - \theta_1 \bar{x}$

- Notations

- x : the independent variable
- y : the dependent variable
- m : the number of training examples
- $x^{(i)}$: the input value of the i^{th} training example
- $y^{(i)}$: the target value of the i^{th} training example
- $\hat{y}^{(i)}$: the prediction made on the i^{th} training example by the current hypothesis function
- α : the learning rate; determines how big steps we take when updating the θ parameters

- Hyperparameters:

- α : if too small, slow gradient descent; if too large, gradient descent may fail to converge

- Problems:

- gradient descent may converge to a local minimum

1.2 Multivariate linear regression

- Description

Multivariate linear regression is a statistical method that studies the relationship between multiple variables:

- n variables $x = \{x_1, x_2, \dots, x_n\}$: the predictor, explanatory, independent variables,
- one y variable: the response, outcome, dependent variable.

- Model's hypothesis

$$h(x) = \theta_0 + \theta_1 * x_1 + \theta_2 * x_2 + \dots + \theta_n * x_n = \hat{y}$$

- Model's parameters (# $n + 1$)

$$\theta = \{\theta_0, \theta_1, \dots, \theta_n\}$$

- Cost function (in this case, the squared error function)

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)})^2$$

- Goal

$$\text{minimize}_{\theta} J(\theta)$$

- Algorithm

- gradient descent
 - start with some initial values for $\theta_0, \theta_1, \dots, \theta_n$ (usually 0)
 - keep changing θ s to reduce $J(\theta)$
 - $\theta_0 = \theta_0 - \alpha \frac{\partial J}{\partial \theta_0} = \theta_0 - \alpha \left(\frac{1}{m} \sum_{i=1}^m (\hat{y}^i - y^{(i)}) \right)$
 - $\theta_1 = \theta_1 - \alpha \frac{\partial J}{\partial \theta_1} = \theta_1 - \alpha \left(\frac{1}{m} \sum_{i=1}^m (\hat{y}^i - y^{(i)}) \cdot x_1^{(i)} \right)$
 - ...
 - $\theta_n = \theta_n - \alpha \frac{\partial J}{\partial \theta_n} = \theta_n - \alpha \left(\frac{1}{m} \sum_{i=1}^m (\hat{y}^i - y^{(i)}) \cdot x_n^{(i)} \right)$

- Notations

- x : the independent variables
- y : the dependent variable
- m : the number of training examples
- n : the number of features representing each training example
- $x^{(i)}$: the input values of the i^{th} training example
- $x_j^{(i)}$: the value of the j^{th} feature of the i^{th} training example

- $y^{(i)}$: the target value of the i^{th} training example
 - $\hat{y}^{(i)}$: the prediction made on the i^{th} training example by the current hypothesis function
 - α : the learning rate; determines how big steps we take when updating the θ parameters
- Hyperparameters:
 - α
- Problems:
 - make sure features are on similar scales; gradient descent may be slow otherwise
 - rescaling: $x' = \frac{x - \min(x)}{\max(x) - \min(x)}$
 - mean normalization: $x' = \frac{x - \text{mean}(x)}{\max(x) - \min(x)}$
 - standardization: $x' = \frac{x - \text{mean}(x)}{\text{std}(x)}$
 - make sure the gradient descent is working correctly
 - plot the value of the cost function J over the number of iterations (# *epochs*)
 - for a sufficiently small α , $J(\theta)$ should decrease on every iteration
 - if α is too small, the gradient descent can be slow to converge
 - if α is too large, $J(\theta)$ may not decrease on every iteration; may not converge
 - try values $\alpha \in \{ \dots, 0.001, 0.01, 0.1, 1, \dots \}$
 - in case of **underfitting**
 - try adding new features
 - e.g. use a polynomial regression

$$\theta_0 + \theta_1 * x \Rightarrow \theta_0 + \theta_1 * x + \theta_2 * x^2 + \theta_3 * x^3$$
 - in case of **overfitting**
 - reduce the number of features
 - manually select which features to keep
 - use a model-selection algorithm
 - use regularization
 - keep all the features, but reduce the magnitude of parameters θ ;
works well when working with a lot of features, each of which contributes a bit to predicting y
 - intuition: *shrink* model parameters in order to *smooth out* the decision boundary (generate a *simpler* hypothesis)
 - cost function with the *regularization term*

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)})^2 + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$
 - parameter update in gradient descent (for $j \in \{1, 2, \dots, n\}$)

$$\theta_j = \theta_j - \alpha \frac{\partial J}{\partial \theta_j} = \theta_j - \alpha \left(\frac{1}{m} \sum_{i=1}^m (\hat{y}^i - y^{(i)}) \cdot x_j^{(i)} + \frac{\lambda}{m} \theta_j \right) = \theta_j \left(1 - \alpha \frac{\lambda}{m} \right) - \alpha \frac{1}{m} \sum_{i=1}^m (\hat{y}^i - y^{(i)}) \cdot x_j^{(i)}$$
 - λ is the *regularization parameter* and needs to be tuned;
it controls the trade-off between the goal of fitting the training data well and the goal of keeping the parameters small