


1 Logistic regression (for classification)

Categorize observations based on quantitative features.
Predict target class or probabilities of target classes.



1.1 Binary classification

- Description

Logistic regression is a statistical method that studies the relationship between multiple variables:

- n variables $x = \{x_1, x_2, \dots, x_n\}$: the predictor, explanatory, independent variables,
- one y variable: the response, outcome, dependent variable.

Logistic regression expands the linear regression model with a *logistic function* to make it suitable for classification. Its dependent variable is therefore categorical instead of numerical.

In case of a binary classification task, its dependent variable takes on one out of two possible values

- $y \in \{0, 1\}$
 - 0 indicates the *negative* class
 - 1 indicates the *positive* class
- Model's hypothesis: output the estimated probability that $y = 1$ on input x
 - $z = \theta_0 + \theta_1 * x_1 + \theta_2 * x_2 + \dots + \theta_n * x_n$
 - $h(x) = P(y = 1 | x, \theta) = \sigma(z) = \frac{1}{1 + e^{-z}} = \hat{y} \quad (0 \leq h(x) \leq 1)$
 - Model's parameters ($n + 1$)
 - $\theta = \{\theta_0, \theta_1, \dots, \theta_n\}$
 - Decision boundary
 - linear
 - non-linear when adding extra higher-order polynomial terms to the features
 - Cost function
$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \left(-y^{(i)} \log \hat{y}^{(i)} - (1 - y^{(i)}) \log (1 - \hat{y}^{(i)}) \right)$$

in other words

 - $J(\theta) = \frac{1}{m} \sum_{i=1}^m \left(-\log \hat{y}^{(i)} \right)$ when $y^{(i)} = 1$
 - $J(\theta) = \frac{1}{m} \sum_{i=1}^m \left(-\log (1 - \hat{y}^{(i)}) \right)$ when $y^{(i)} = 0$
 - Goal
$$\text{minimize}_{\theta} J(\theta)$$
 - Algorithm
 - gradient descent
 - start with some initial values for $\theta_0, \theta_1, \dots, \theta_n$ (usually normal random values)
 - keep changing θ s to reduce $J(\theta)$

- $\theta_0 = \theta_0 - \alpha \frac{\partial J}{\partial \theta_0} = \theta_0 - \alpha \left(\frac{1}{m} \sum_{i=1}^m (\hat{y}^i - y^{(i)}) \right)$
- $\theta_1 = \theta_1 - \alpha \frac{\partial J}{\partial \theta_1} = \theta_1 - \alpha \left(\frac{1}{m} \sum_{i=1}^m (\hat{y}^i - y^{(i)}) \cdot x_1^{(i)} \right)$
- ...
- $\theta_n = \theta_n - \alpha \frac{\partial J}{\partial \theta_n} = \theta_n - \alpha \left(\frac{1}{m} \sum_{i=1}^m (\hat{y}^i - y^{(i)}) \cdot x_n^{(i)} \right)$

- Hyperparameters

- α

- Problems

- the idea of feature scaling also applied for logistic regression
- make sure the gradient descent is working correctly
- in case of **underfitting**

- try adding new features
 - e.g. use a polynomial regression

$$\theta_0 + \theta_1 * x \Rightarrow \theta_0 + \theta_1 * x + \theta_2 * x^2 + \theta_3 * x^3$$

- in case of **overfitting**

- reduce the number of features
 - manually select which features to keep
 - use a model-selection algorithm
- use regularization
 - keep all the features, but reduce the magnitude of parameters θ ;
works well when working with a lot of features, each of which contributes a bit to predicting y
 - intuition: *shrink* model parameters in order to *smooth out* the decision boundary (generate a *simpler* hypothesis)
 - cost function with the regularization term

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \left(-y^{(i)} \log \hat{y}^{(i)} - (1 - y^{(i)}) \log (1 - \hat{y}^{(i)}) \right) + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

- parameter update in gradient descent (for $j \in \{1, 2, \dots, n\}$)

$$\theta_j = \theta_j - \alpha \frac{\partial J}{\partial \theta_j} = \theta_j - \alpha \left(\frac{1}{m} \sum_{i=1}^m (\hat{y}^i - y^{(i)}) \cdot x_j^{(i)} + \frac{\lambda}{m} \theta_j \right) = \theta_j \left(1 - \alpha \frac{\lambda}{m} \right) - \alpha \frac{1}{m} \sum_{i=1}^m (\hat{y}^i - y^{(i)}) \cdot x_j^{(i)}$$

- λ is the *regularization parameter* and needs to be tuned;
it controls the trade-off between the goal of fitting the training data well and the goal of keeping the parameters small

- Evaluation metrics

- accuracy: out of all predictions, how many of them were correct
- precision: out of all *positive* predictions, how many of them were actually *positive* examples
- recall: out of all positive examples, how many of them were detected as positive
- f1-score: harmonic mean between precision and recall

- Trading off precision and recall

- use a different threshold for making decisions whether the class is positive or negative
- minimize the number of false positives
 - generate a higher precision, but a lower recall
 - choose a higher threshold (e.g. $\sigma = \{0.7, \dots, 0.9\}$)
- minimize the number of false negatives
 - generate a higher recall, but a lower precision
 - choose a lower threshold (e.g. $\sigma = \{0.1, \dots, 0.3\}$)
- plot the precision-recall tradeoff curve by testing various threshold values between $[0, 1]$

1.2 Multiclass classification

- Description
 - use the one-vs-all (one-vs-rest) approach
 - turn the problem into C binary classification problems (generate C decision boundaries)
 - formally: train a logistic regression classifier $h^{(i)}(x)$ for each class i to predict the probability that $y = i$
 - on a new input x , in order to make a prediction pick the class i that maximizes $\max_i h^{(i)}(x)$